

Offline Semantic Search Browser Extension: A Privacy-First, High-Speed Architecture for Meaning-Based Document Retrieval Using Quantized AI Embeddings

Varun Attri, Md Sami Khan, Tanmay Singh

Department of Computer Science and Engineering

Noida Institute of Engineering and Technology

Dr. A.P.J. Abdul Kalam Technical University, Uttar Pradesh,
India

Email: mdsamik9@gmail.com

Abstract—This paper presents the design and implementation of a fully offline, privacy-first semantic search browser extension built on Manifest V3 for Chromium-based browsers. Unlike conventional keyword-based search tools, the proposed system enables meaning-based retrieval across webpages, PDF documents, and local text files, entirely within the user's browser—without any cloud API calls after the initial model download. The architecture integrates a quantized INT8 ONNX version of the all-MiniLM-L6-v2 sentence-transformer model (approximately 22 MB), a Web Worker-based inference pipeline, IndexedDB persistent vector storage via Dexie.js, and cosine similarity search augmented with HNSW indexing. A layered Chrome Extension architecture comprising a service worker message router, offscreen document for AI inference, content script for DOM extraction, and a React 18 popup UI is described in detail. The system achieves sub-1-second query latency, under-3-second page indexing, and a memory footprint below 200 MB. Experimental evaluation confirms that the design satisfies practical performance targets while providing strong privacy guarantees. This work contributes a reproducible blueprint for deploying transformer-based semantic search at the browser edge, with potential applications in personal knowledge management, offline research assistance, and privacy-sensitive enterprise environments.

Keywords—semantic search; browser extension; ONNX Runtime Web; sentence transformers; vector embeddings; IndexedDB; on-device AI; privacy-first; Manifest V3; HNSW; all-MiniLM-L6-v2; offline NLP

I. INTRODUCTION

Modern web browsers have evolved into the primary interface through which users consume, annotate, and retrieve information. Despite this centrality, the search capabilities available to users within a browsing session remain predominantly lexical: exact keyword matching or simple substring search. These approaches fail to capture semantic relationships—a user searching for "vehicle engine problem" will not retrieve documents discussing "car motor malfunction," even though the concepts are equivalent. Neural information retrieval, powered by dense vector representations

of text, has demonstrated significant improvements over term-based methods [1]. However, the deployment of such models has been largely confined to server-side cloud infrastructure, introducing latency, privacy risks, and connectivity requirements.

This work bridges that gap by demonstrating that a production-quality semantic search system can be deployed entirely inside a web browser as a Chrome extension. The key contributions of this paper are as follows:

- A complete Manifest V3 browser extension architecture for offline AI inference, including a service worker message router, offscreen document inference host, IIFE content script, and React 18 popup.
- Integration of a quantized (INT8) ONNX sentence-transformer model delivering 384-dimensional embeddings with a disk footprint of approximately 22 MB.
- A Web Worker embedding pipeline with IndexedDB-backed vector storage and cosine similarity search, achieving sub-1-second retrieval latency over thousands of indexed chunks.
- A Content Security Policy-compliant build system (Vite 6 / Rollup) that rewrites CDN references to local WASM assets and packages ONNX Runtime files for offline execution.
- Quantitative evaluation of latency, memory, and storage against defined performance targets.

The remainder of this paper is organized as follows. Section II reviews related work. Section III describes the system architecture. Section IV details the core engine modules. Section V presents the build and deployment design. Section VI reports experimental evaluation. Section VII discusses privacy and security properties. Section VIII concludes with future work directions.

II. RELATED WORK

A. Dense Passage Retrieval and Semantic Embeddings

The shift from sparse term-based retrieval (BM25, TF-IDF) to dense vector retrieval was catalyzed by the bi-encoder model DPR [2] and further advanced by sentence-transformers such as SBERT [3]. All-MiniLM-L6-v2 [4] offers a highly compact yet accurate sentence embedding model trained via knowledge distillation, making it suitable for resource-constrained deployments. Our work applies INT8 quantization via ONNX Runtime [5] to further reduce the model footprint to approximately 22 MB without material accuracy degradation.

B. On-Device and Edge AI Inference

WebAssembly (Wasm) and the ONNX Runtime Web backend have enabled neural network inference directly in the browser sandbox [6]. The @huggingface/transformers JavaScript library abstracts model loading and tokenization, wrapping ONNX Runtime Web with a pipeline API. Prior work on in-browser ML has covered image classification [7] and speech recognition [8], but semantic text retrieval in a browser extension context remains underexplored. Our architecture extends this body of work with a multi-process Chrome extension design that satisfies Manifest V3 constraints.

C. Browser Extension Search Tools

Existing browser search extensions such as Ctrl+F enhancements, ReadAloud, and cloud-backed AI extensions (e.g., Perplexity, Monica) either operate on exact-match principles or transmit content to remote servers. The privacy implications of cloud-mediated search have been studied in [9]. Our system provides an entirely local alternative, with no network egress after model download, aligning with privacy-by-design principles [10].

D. Approximate Nearest Neighbour Search

For large corpora, exhaustive cosine similarity search scales as $O(n)$ in the number of stored vectors. Hierarchical Navigable Small World (HNSW) graphs [11] reduce this to approximate $O(\log n)$ queries. The proposed system implements cosine similarity as the primary search path for typical browsing corpora (hundreds to a few thousand chunks), with HNSW available as an optional index for larger collections.

III. SYSTEM ARCHITECTURE

The extension is structured as a set of isolated processes mandated by the Chrome Manifest V3 model. Fig. 1 illustrates the four primary process contexts and their communication topology.




Fig. 1. Extension process topology and inter-process communication.

A. Service Worker (Message Router)

The service worker (src/background/service-worker.ts) acts as the central hub for all inter-process message passing. It routes indexing requests from the popup or content script to the offscreen document and returns results to the requester. Because Manifest V3 service workers are ephemeral, all persistent state is stored in IndexedDB rather than in-memory variables.

B. Content Script (DOM Text Extractor)

The content script (src/content/content-script.ts) is injected into every web page as an IIFE bundle. Chrome content scripts cannot use ES modules; the Vite build therefore compiles this entry point separately in IIFE format. The script extracts the visible text content from the DOM, strips scripts and styles, and streams the cleaned text to the service worker on request.

C. Offscreen Document (AI Inference Host)

Chrome Manifest V3 prohibits direct use of Web Workers from the service worker context. The offscreen document (src/offscreen/) provides a persistent DOM context in which a dedicated embedding Web Worker can be instantiated. The offscreen document receives tokenization and inference requests, delegates them to the embedding worker, and returns the resulting 384-dimensional float vectors.

D. Popup UI (React 18 / Tailwind CSS)

The popup window (src/popup/) is a React 18 single-page application styled with Tailwind CSS using a custom dark-

mode palette built around violet-purple accent tokens. It provides index management, query input, result display with confidence scores and excerpt highlighting, and settings (including opt-in auto-indexing on page visit).




Fig. 2. Flowchart

IV. CORE ENGINE MODULES

A. Text Processor

The text processor (src/core/text-processor.ts) accepts raw HTML or plain text and produces an array of semantically coherent chunks of 200–500 tokens. Processing stages include: (1) HTML tag and script/style stripping, (2) Unicode normalization and whitespace collapsing, (3) sentence boundary detection using regex heuristics, and (4) greedy sentence packing into chunks with configurable overlap to preserve cross-boundary context.

B. Embedding Engine

The embedding engine (src/core/embedding-engine.ts) wraps the @huggingface/transformers pipeline API targeting the Xenova/all-MiniLM-L6-v2 ONNX quantized checkpoint. The model is loaded once and cached; subsequent calls reuse the loaded session. The engine runs inside the embedding Web Worker (src/workers/embedding-worker.ts), keeping the main thread non-blocking. ONNX Runtime WASM files (ort-wasm-simd-threaded.*) are served from the extension's local /wasm/

directory to comply with the extension's Content Security Policy.

C. Similarity Search

The similarity search module (`src/core/similarity-search.ts`) implements mean-pooled cosine similarity between the query embedding and all stored chunk embeddings. Scores are ranked in descending order, and the top-k results (default $k=10$) are returned with their source URL, chunk text, and score. An optional HNSW index is constructed over the stored embedding matrix when corpus size exceeds a configurable threshold, reducing retrieval complexity for large collections.

D. PDF Extractor

The PDF extractor (`src/core/pdf-extractor.ts`) leverages `pdfjs-dist 4.x` with its worker served locally from `dist/pdf/pdf.worker.min.mjs`. It extracts text page-by-page, concatenates the output, and passes it to the text processor. This enables the same semantic indexing and search pipeline to operate over uploaded PDF documents without any server-side component.

E. Cache Manager and IndexedDB Schema

All indexed data is persisted in IndexedDB via `Dexie.js` (`src/db/`). The schema stores: (1) a page metadata record keyed by URL with timestamp and content hash, (2) an array of chunk records each containing the chunk text and its serialized embedding `Float32Array`, and (3) HNSW graph nodes when the optional index is active. Cache invalidation compares the stored content hash against the current page hash on re-visit, triggering re-indexing only when content changes.

V. BUILD AND DEPLOYMENT DESIGN

A. Multi-Entry Vite 6 Build

The build system uses Vite 6 with Rollup to produce four separate entry points: (1) popup HTML, (2) the service worker as an ES module bundle, (3) the offscreen document HTML, and (4) the embedding Web Worker. The content script is built as a fifth separate bundle in IIFE format via a custom Vite plugin (`buildContentScriptPlugin`), because Chrome content scripts cannot execute ES modules.

B. WASM Asset Packaging

A custom `localWasmPlugin` copies four ONNX Runtime WASM and MJS files from `node_modules/onnxruntime-web/dist` to `dist/wasm/` at bundle time. A companion `rewriteOnnxCdnUrls` plugin rewrites any residual `jsdelivr` CDN URLs in the bundle to the local `/wasm/` path, ensuring that the extension's Content Security Policy (`script-src 'self' 'wasm-unsafe-eval'`) is never violated.

C. Content Security Policy

The manifest specifies a strict extension_pages CSP: `script-src 'self' 'wasm-unsafe-eval'; object-src 'self'`. The `'wasm-unsafe-eval'` directive is the minimum required to execute compiled ONNX WASM modules. No external script sources are permitted. All model weights, WASM binaries, and worker scripts are bundled into the extension package.

D. TypeScript Configuration

The project targets ES2022 with strict mode, `moduleResolution` set to `bundler`, and path aliases (`@/`, `@core/`, `@db/`, `@utils/`, `@types/`) mapped to the `src/` subtree. Type checking is performed as a pre-build step (`tsc --noEmit`) and enforced in CI via the build script defined in `package.json`.

VI. EXPERIMENTAL EVALUATION

Performance was measured on a mid-range laptop (Intel Core i5-1135G7, 8 GB RAM, Chrome 124) to establish representative real-world baselines. Table I summarizes the achieved results against the design targets.

TABLE I
PERFORMANCE METRICS: DESIGN TARGETS VS. ACHIEVED RESULTS

Metric	Target	Achieved
Search latency (1000 chunks)	< 1 second	~180 ms
Page indexing time	< 3 seconds	~1.4 s (avg)
Model load (first run)	< 5 seconds	~3.2 s
Peak memory usage	< 200 MB	~145 MB
Storage per page	50–200 KB	~85 KB (avg)
Model disk footprint	—	~22 MB (INT8)

Search latency of approximately 180 ms for 1,000 stored chunks is well within the 1-second target. The dominant cost is the embedding of the query string (~95 ms) rather than the vector comparison (~85 ms). For very large corpora (> 5,000 chunks), enabling the HNSW index reduces retrieval time to approximately 40 ms at the cost of an approximately 12 MB index structure in memory.

Model loading exhibits a cold-start latency of approximately 3.2 seconds on the first extension activation, after which the ONNX session is cached in the Web Worker and reused for the lifetime of the offscreen document. Subsequent activations within the same browser session complete in under 200 ms.

Memory consumption peaks at approximately 145 MB during simultaneous model inference and IndexedDB batch write, comfortably below the 200 MB budget. Idle consumption after indexing stabilizes at approximately 55 MB.

VII. PRIVACY AND SECURITY ANALYSIS

The system is designed to satisfy privacy-by-design principles [10]. Table II summarizes key properties.

TABLE II
PRIVACY AND SECURITY PROPERTIES

Property	Implementation
Data locality	All vectors and text stored exclusively in browser IndexedDB; no external database.
Network isolation	Zero outbound requests after model download; no telemetry or analytics endpoints.
Minimal permissions	Only <code>activeTab</code> , <code>storage</code> , <code>offscreen</code> , <code>sidePanel</code> , <code>scripting</code> ; no identity or cookie permissions.

CSP enforcement	script-src 'self' 'wasm-unsafe-eval'; object-src 'self' in manifest.
Auto-index opt-in	Auto-indexing on page visit is disabled by default; user must explicitly enable.
Content integrity	Model weights served from extension package; no runtime CDN fetch.

The extension requests host_permissions over <all_urls> solely to enable the content script injection required for DOM text extraction. No page content, query strings, or embedding vectors are transmitted outside the browser process. The IndexedDB storage is sandboxed per browser profile and is not accessible to web page JavaScript.

VIII. CONCLUSION AND FUTURE WORK

This paper has presented the design, implementation, and evaluation of a fully offline, privacy-first semantic search browser extension. By combining a quantized ONNX sentence-transformer model, a multi-process Chrome extension architecture, and a Vite-based build system that packages all WASM and model assets locally, the system achieves sub-1-second semantic query latency with no cloud dependency after initial setup. The privacy-by-design architecture ensures that no user content leaves the browser, making it suitable for privacy-sensitive deployments in healthcare, legal, and enterprise research contexts.

Several directions for future work are identified. First, cross-browser support for Firefox (Manifest V2/V3) and Safari (Web Extensions API) would broaden reach. Second, the adoption of a pure JavaScript HNSW library would eliminate the remaining $O(n)$ fallback for large corpora without additional WASM dependencies. Third, incremental re-indexing—tracking DOM mutations via MutationObserver—would ensure that dynamically updated pages remain current without full re-indexing. Fourth, keyboard navigation of results (arrow keys + Enter) and a persistent search history dropdown are planned UI enhancements. Finally, federated search across multiple devices using the Chrome Storage Sync API, with end-to-end encrypted embedding vectors, is an open research challenge that would extend the personal knowledge management use case.

. REFERENCES

- [1] J. Lin and X. Ma, "A Few Brief Notes on DeepImpact, COIL, and a Conceptual Framework for Information Retrieval Techniques," arXiv preprint arXiv:2106.14807, 2021.
- [2] V. Karpukhin et al., "Dense Passage Retrieval for Open-Domain Question Answering," in Proc. EMNLP, 2020, pp. 6769–6781.
- [3] N. Reimers and I. Gurevych, "Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks," in Proc. EMNLP, 2019, pp. 3982–3992.
- [4] M. Wang et al., "MiniLM: Deep Self-Attention Distillation for Task-Agnostic Compression of Pre-Trained Transformers," in Advances in Neural Information Processing Systems (NeurIPS), 2020, vol. 33, pp. 5776–5788.
- [5] ONNX Runtime Team, "ONNX Runtime: Cross-Platform, High-Performance ML Inferencing," Microsoft Open Source, 2023. [Online]. Available: <https://onnxruntime.ai>

- [6] A. Haas et al., "Bringing the Web up to Speed with WebAssembly," in Proc. ACM PLDI, 2017, pp. 185–200.
- [7] J. Herrmann and T. Mitchell, "TensorFlow.js: Machine Learning for the Web and Beyond," in Proc. SysML, 2019.
- [8] A. Baevski et al., "wav2vec 2.0: A Framework for Self-Supervised Learning of Speech Representations," in Advances in NeurIPS, 2020, vol. 33.
- [9] C. Castelluccia and M. Naini, "On the Privacy Implications of Browser Extensions," in Proc. IEEE S&P Workshops, 2020, pp. 1–6.
- [10] A. Cavoukian, "Privacy by Design: The 7 Foundational Principles," Office of the Information and Privacy Commissioner of Ontario, 2009.
- [11] Y. Malkov and D. Yashunin, "Efficient and Robust Approximate Nearest Neighbor Search Using Hierarchical Navigable Small World Graphs," IEEE Trans. Pattern Anal. Mach. Intell., vol. 42, no. 4, pp. 824–836, 2020.